

Assignment-04-Simple Linear Regression-1

```
In [5]: # import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
```

```
In [6]: # import dataset
dataset=pd.read_csv('delivery_time.csv')
dataset
```

```
Out[6]:
```

	Delivery Time	Sorting Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10
5	15.35	6
6	19.00	7
7	9.50	3
8	17.90	10
9	18.75	9
10	19.83	8
11	10.75	4
12	16.68	7
13	11.50	3
14	12.03	3
15	14.88	4
16	13.75	6
17	18.11	7
18	8.00	2
19	17.83	7
20	21.50	5

EDA and Data Visualization

```
In [9]: dataset=dataset.rename(columns={'Delivery Time': 'dt','Sorting Time': 'st' })
```


```
In [30]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    dt      21 non-null      float64
 1    st      21 non-null      int64
dtypes: float64(1), int64(1)
memory usage: 464.0 bytes
```

```
In [32]: sns.distplot(dataset['dt'])
```

```
C:\Users\HP\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```


```
Out[32]: <AxesSubplot:xlabel='dt', ylabel='Density'>
```



```
In [33]: sns.distplot(dataset['st'])
```

```
C:\Users\HP\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[33]: <AxesSubplot:xlabel='st', ylabel='Density'>
```



Feature Engineering

```
In [34]: dataset=dataset.rename(columns={'Delivery Time': 'dt','Sorting Time': 'st' })
```

Correlation Analysis


```
In [10]: dataset.corr()
```

```
Out[10]:
```

	dt	st
dt	1.000000	0.825997
st	0.825997	1.000000


```
In [11]: plt.scatter(x=dataset.st, y=dataset.dt, color='green')  
plt.xlabel("Sorting time")  
plt.ylabel("Delivery time")
```

```
Out[11]: Text(0, 0.5, 'Delivery time')
```




```
In [12]: plt.boxplot(dataset.dt)
```

```
Out[12]: {'whiskers': [<matplotlib.lines.Line2D at 0x19226575640>,  
<matplotlib.lines.Line2D at 0x19226575940>],  
'caps': [<matplotlib.lines.Line2D at 0x19226575cd0>,  
<matplotlib.lines.Line2D at 0x19226575ee0>],  
'boxes': [<matplotlib.lines.Line2D at 0x19226575370>],  
'medians': [<matplotlib.lines.Line2D at 0x192265881f0>],  
'fliers': [<matplotlib.lines.Line2D at 0x192265884c0>],  
'means': []}
```



```
In [13]: plt.hist(dataset.dt, bins=5)
```

```
Out[13]: (array([5., 4., 8., 3., 1.]),
          array([ 8. , 12.2, 16.4, 20.6, 24.8, 29. ]),
          <BarContainer object of 5 artists>)
```



Model Building

```
In [14]: model12=smf.ols("dt~st",data=dataset).fit()
```

Model Testing

```
In [35]: # Finding Coefficient parameters
          model12.params
```

```
Out[35]: Intercept    6.582734
          st          1.649020
          dtype: float64
```

```
In [16]: model12.summary()
```

Out[16]:

OLS Regression Results

Dep. Variable:		dt	R-squared:		0.682	
Model:		OLS	Adj. R-squared:		0.666	
Method:		Least Squares		F-statistic:		40.80
Date:		Tue, 29 Nov 2022		Prob (F-statistic):		3.98e-06
Time:		13:48:20		Log-Likelihood:		-51.357
No. Observations:		21		AIC:		106.7
Df Residuals:		19		BIC:		108.8
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.5827	1.722	3.823	0.001	2.979	10.186
st	1.6490	0.258	6.387	0.000	1.109	2.189
Omnibus:		3.649	Durbin-Watson:		1.248	
Prob(Omnibus):		0.161	Jarque-Bera (JB):		2.086	
Skew:		0.750	Prob(JB):		0.352	
Kurtosis:		3.367	Cond. No.		18.3	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [18]: model3=smf.ols("dt~np.log(st)",data=dataset).fit()

In [19]: model3.params

Out[19]: Intercept      1.159684
np.log(st)      9.043413
dtype: float64

In [20]: model3.summary()
```

Out[20]:

OLS Regression Results

Dep. Variable:	dt	R-squared:	0.695
Model:	OLS	Adj. R-squared:	0.679
Method:	Least Squares	F-statistic:	43.39
Date:	Tue, 29 Nov 2022	Prob (F-statistic):	2.64e-06
Time:	13:51:47	Log-Likelihood:	-50.912
No. Observations:	21	AIC:	105.8
Df Residuals:	19	BIC:	107.9
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.1597	2.455	0.472	0.642	-3.978	6.297
np.log(st)	9.0434	1.373	6.587	0.000	6.170	11.917

Omnibus:	5.552	Durbin-Watson:	1.427
Prob(Omnibus):	0.062	Jarque-Bera (JB):	3.481
Skew:	0.946	Prob(JB):	0.175
Kurtosis:	3.628	Cond. No.	9.08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [21]: model2.conf_int(0.05) # 95% confidence interval
```

```
Out[21]:
```

	0	1
Intercept	2.979134	10.186334
st	1.108673	2.189367

```
In [22]: model3.conf_int(0.05) # 95% confidence interval
```

```
Out[22]:
```

	0	1
Intercept	-3.97778	6.297147
np.log(st)	6.16977	11.917057

Model Predictions


```
In [24]: pred2 = model2.predict(dataset) # Predicted values of dt using the model
```

```
In [25]: pred3 = model3.predict(dataset) # Predicted values of dt using the model
```

```
In [26]: plt.scatter(x=dataset.st, y=dataset.dt, color='green')
plt.plot(dataset.st, pred2, color='black')
```


```
plt.xlabel("Sorting time")
plt.ylabel("Delivery time")
```

Out[26]: Text(0, 0.5, 'Delivery time')



```
In [29]: plt.scatter(x=dataset.st, y=dataset.dt, color='red')
plt.plot(dataset.st, pred3,color='green')
plt.xlabel("Sorting time")
plt.ylabel("Delivery time")
```

Out[29]: Text(0, 0.5, 'Delivery time')



Model Testing

```
In [70]: # Finding Coefficient parameters
model.params
```

Out[70]: Intercept 25792.200199
YearsExperience 9449.962321
dtype: float64

```
In [71]: # Finding tvalues and pvalues
model.tvalues , model.pvalues
```

```
Out[71]: (Intercept          11.346940
          YearsExperience    24.950094
          dtype: float64,
          Intercept          5.511950e-12
          YearsExperience     1.143068e-20
          dtype: float64)
```

```
In [72]: # Finding Rsquared Values
         model.rsquared , model.rsquared_adj
```

```
Out[72]: (0.9569566641435086, 0.9554194021486339)
```

Assignment-04-Simple Linear Regression-2

```
In [36]: import pandas as pd
         import numpy as np
         import scipy.stats as stats
         import matplotlib.pyplot as plt
         import seaborn as sns
         import statsmodels.api as smf
         import statsmodels.formula.api as sm
         import warnings
         warnings.filterwarnings('ignore')
```

```
In [37]: df = pd.read_csv('Salary_Data.csv')
         df
```


Out[37]:

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

In [38]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

In [39]: `df.describe()`

Loading [MathJax]/extensions/Safe.js

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

Checking for Null Values

```
In [40]: df.isnull().sum()
```

```
Out[40]: YearsExperience    0
Salary                  0
dtype: int64
```

Checking for Duplicate Values

```
In [41]: df[df.duplicated()].shape
```


```
Out[41]: (0, 2)
```

```
In [42]: df[df.duplicated()]
```

```
Out[42]:   YearsExperience  Salary
```

Plotting the data to check for outliers

```
In [43]: plt.subplots(figsize = (9,6))
plt.subplot(121)
plt.boxplot(df['Salary'])
plt.title('Salary Hike')
plt.subplot(122)
plt.boxplot(df['YearsExperience'])
plt.title('Years of Experience')
plt.show()
```



Checking the Correlation between variables

In [44]: `df.corr()`

Out[44]:


	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

Visualization of Correlation between x and y

regplot = regression plot


In [45]: `sns.regplot(x=df['YearsExperience'], y=df['Salary'])`

Out[45]: `<AxesSubplot:xlabel='YearsExperience', ylabel='Salary'>`



Checking for Homoscedasticity or Heteroscedasticity

```
In [46]: plt.figure(figsize = (8,6), facecolor = 'lightblue')
sns.scatterplot(x = df['YearsExperience'], y = df['Salary'])
plt.title('Homoscedasticity', fontweight = 'bold', fontsize = 16)
plt.show()
```


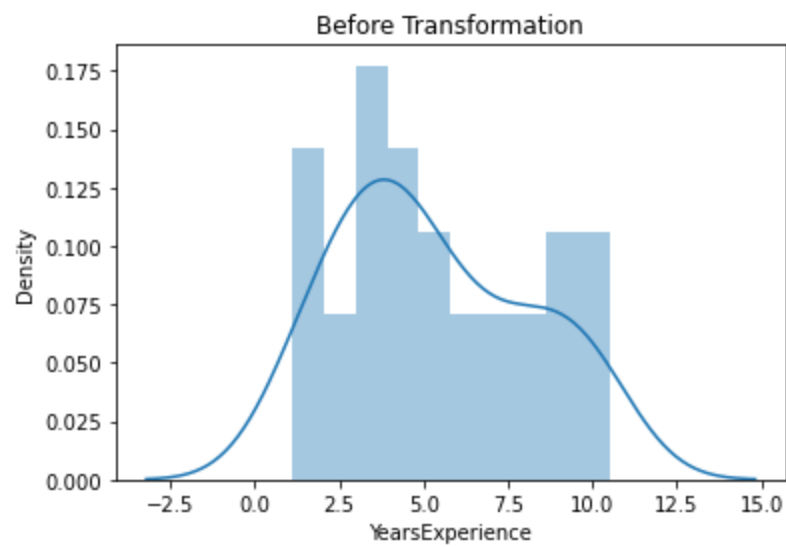


```
In [47]: df.var()
```


```
Out[47]: YearsExperience    8.053609e+00
Salary                  7.515510e+08
dtype: float64
```

Feature Engineering

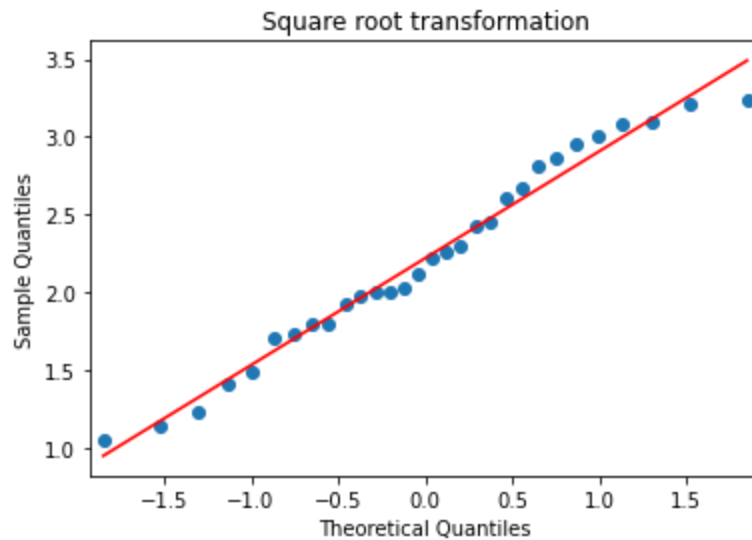


```
In [48]: sns.distplot(df['YearsExperience'], bins = 10, kde = True)
plt.title('Before Transformation')
sns.distplot(np.log(df['YearsExperience']), bins = 10, kde = True)
plt.title('After Transformation')
plt.show()
```




```
In [49]: labels = ['Before Transformation', 'After Transformation']
sns.distplot(df['YearsExperience'], bins = 10, kde = True)
sns.distplot(np.log(df['YearsExperience']), bins = 10, kde = True)
plt.legend(labels)
plt.show()
```





```
In [50]: smf.qqplot(np.log(df['YearsExperience']), line = 'r')
plt.title('No transformation')
smf.qqplot(np.sqrt(df['YearsExperience']), line = 'r')
plt.title('Log transformation')
smf.qqplot(np.sqrt(df['YearsExperience']), line = 'r')
plt.title('Square root transformation')
smf.qqplot(np.cbrt(df['YearsExperience']), line = 'r')
plt.title('Cube root transformation')
plt.show()
```


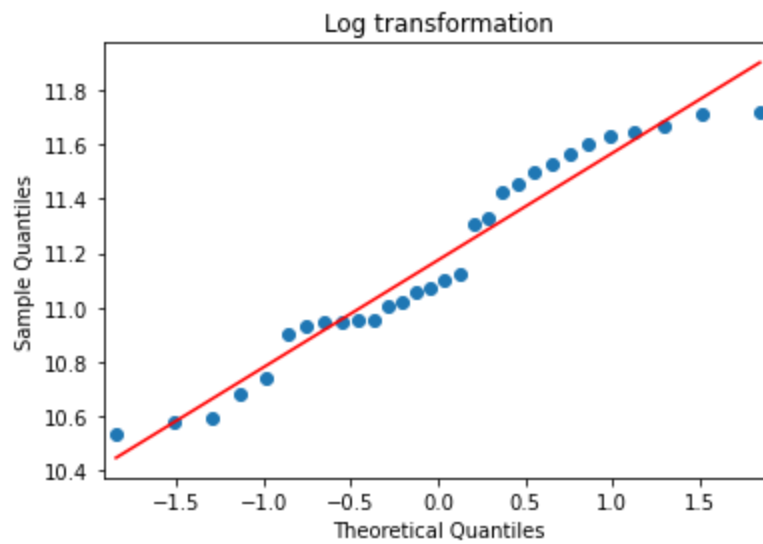






```
In [51]: labels = ['Before Transformation', 'After Transformation']
sns.distplot(df['Salary'], bins = 10, kde = True)
sns.distplot(np.log(df['Salary']), bins = 10, kde = True)
plt.title('After Transformation')
plt.show()
```



```
In [52]: smf.qqplot(df['Salary'], line = 'r')
plt.title('No transformation')
smf.qqplot(np.log(df['Salary']), line = 'r')
plt.title('Log transformation')
smf.qqplot(np.sqrt(df['Salary']), line = 'r')
plt.title('Square root transformation')
smf.qqplot(np.cbrt(df['Salary']), line = 'r')
plt.title('Cube root transformation')
plt.show()
```





Fitting a Linear Regression Model

```
In [53]: import statsmodels.formula.api as sm
model = sm.ols('Salary~YearsExperience', data = df).fit()
```

```
In [54]: model.summary()
```

```
Out[54]:
```

OLS Regression Results						
Dep. Variable:	Salary	R-squared:	0.957			
Model:	OLS	Adj. R-squared:	0.955			
Method:	Least Squares	F-statistic:	622.5			
Date:	Tue, 29 Nov 2022	Prob (F-statistic):	1.14e-20			
Time:	15:37:31	Log-Likelihood:	-301.44			
No. Observations:	30	AIC:	606.9			
Df Residuals:	28	BIC:	609.7			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.579e+04	2273.053	11.347	0.000	2.11e+04	3.04e+04
YearsExperience	9449.9623	378.755	24.950	0.000	8674.119	1.02e+04
Omnibus:	2.140	Durbin-Watson:	1.648			
Prob(Omnibus):	0.343	Jarque-Bera (JB):	1.569			
Skew:	0.363	Prob(JB):	0.456			
Kurtosis:	2.147	Cond. No.	13.2			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [55]: model1 = sm.ols('np.sqrt(Salary)~np.sqrt(YearsExperience)', data = df).fit()
```

```
model11.summary()
```

Out[55]:

OLS Regression Results

Dep. Variable:	np.sqrt(Salary)	R-squared:	0.942
Model:	OLS	Adj. R-squared:	0.940
Method:	Least Squares	F-statistic:	454.3
Date:	Tue, 29 Nov 2022	Prob (F-statistic):	7.58e-19
Time:	15:38:01	Log-Likelihood:	-116.52
No. Observations:	30	AIC:	237.0
Df Residuals:	28	BIC:	239.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	103.5680	8.178	12.663	0.000	86.815	120.321
np.sqrt(YearsExperience)	75.6269	3.548	21.315	0.000	68.359	82.895

Omnibus:	0.924	Durbin-Watson:	1.362
Prob(Omnibus):	0.630	Jarque-Bera (JB):	0.801
Skew:	0.087	Prob(JB):	0.670
Kurtosis:	2.219	Cond. No.	9.97

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [56]: model12 = sm.ols('np.cbrt(Salary)~np.cbrt(YearsExperience)', data = df).fit()  
model12.summary()
```

Out [56]:

OLS Regression Results			
Dep. Variable:	np.cbrt(Salary)	R-squared:	0.932
Model:	OLS	Adj. R-squared:	0.930
Method:	Least Squares	F-statistic:	386.5
Date:	Tue, 29 Nov 2022	Prob (F-statistic):	6.37e-18
Time:	15:38:23	Log-Likelihood:	-50.589
No. Observations:	30	AIC:	105.2
Df Residuals:	28	BIC:	108.0
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	16.6603	1.300	12.811	0.000	13.996	19.324
np.cbrt(YearsExperience)	14.8963	0.758	19.659	0.000	13.344	16.448

Omnibus:	0.386	Durbin-Watson:	1.229
Prob(Omnibus):	0.824	Jarque-Bera (JB):	0.535
Skew:	0.070	Prob(JB):	0.765
Kurtosis:	2.361	Cond. No.	12.0

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [57]: model3 = sm.ols('np.log(Salary)~np.log(YearsExperience)', data = df).fit()  
model3.summary()
```

Out[57]:

OLS Regression Results

Dep. Variable:	np.log(Salary)		R-squared:	0.905		
Model:	OLS		Adj. R-squared:	0.902		
Method:	Least Squares		F-statistic:	267.4		
Date:	Tue, 29 Nov 2022	Prob (F-statistic):	7.40e-16			
Time:	15:38:44		Log-Likelihood:	23.209		
No. Observations:	30		AIC:	-42.42		
Df Residuals:	28		BIC:	-39.61		
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	10.3280	0.056	184.868	0.000	10.214	10.442
np.log(YearsExperience)	0.5621	0.034	16.353	0.000	0.492	0.632
Omnibus:	0.102	Durbin-Watson:	0.988			
Prob(Omnibus):	0.950	Jarque-Bera (JB):	0.297			
Skew:	0.093	Prob(JB):	0.862			
Kurtosis:	2.549	Cond. No.	5.76			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model Testing

```
In [58]: model.params
Out[58]: Intercept      25792.200199
YearsExperience      9449.962321
dtype: float64


In [59]: print(model.tvalues, '\n', model.pvalues)
Intercept      11.346940
YearsExperience  24.950094
dtype: float64
Intercept      5.511950e-12
YearsExperience  1.143068e-20
dtype: float64

In [60]: model.rsquared, model.rsquared_adj
Out[60]: (0.9569566641435086, 0.9554194021486339)
```

Residual Analysis


```
In [61]: import statsmodels.api as sm
Loading [MathJax]/extensions/Safe.js del.resid, line = 'q')
```

```
plt.title('Normal Q-Q plot of residuals of Model without any data transformation')
plt.show()
```



```
In [62]: def get_standardized_values( vals ):
          return (vals - vals.mean())/vals.std()
```

```
In [63]: plt.scatter(get_standardized_values(model.fittedvalues), get_standardized_values(model.r
plt.title('Residual Plot for Model without any data transformation')
plt.xlabel('Standardized Fitted Values')
plt.ylabel('Standardized Residual Values')
plt.show()
```



Model Validation

```
In [64]: from sklearn.metrics import mean_squared_error
```

```
In [65]: model1_pred_y =np.square(model1.predict(df['YearsExperience']))
model2_pred_y =pow(model2.predict(df['YearsExperience']),3)
model3_pred_y =np.exp(model3.predict(df['YearsExperience']))
```

```
In [66]: model1_rmse =np.sqrt(mean_squared_error(df['Salary'], model1_pred_y))
model2_rmse =np.sqrt(mean_squared_error(df['Salary'], model2_pred_y))
model3_rmse =np.sqrt(mean_squared_error(df['Salary'], model3_pred_y))
print('model=', np.sqrt(model.mse_resid),'\n' 'model1=', model1_rmse,'\n' 'model2=', mod
```

```
model= 5788.315051119395  
model1= 5960.64709617431  
model2= 6232.815455835842  
model3= 7219.716974372806
```

```
In [67]: rmse = {'model': np.sqrt(model.mse_resid), 'model1': model1_rmse, 'model2': model3_rmse,  
min(rmse, key=rmse.get)
```

```
Out[67]: 'model'
```

Predicting values

```
In [69]: # first model results without any transformation  
predicted2 = pd.DataFrame()  
predicted2['YearsExperience'] = df.YearsExperience  
predicted2['Salary'] = df.Salary  
predicted2['Predicted_Salary_Hike'] = pd.DataFrame(model.predict(predicted2.YearsExperience  
predicted2
```

Out[69]:

	YearsExperience	Salary	Predicted_Salary_Hike
0	1.1	39343.0	36187.158752
1	1.3	46205.0	38077.151217
2	1.5	37731.0	39967.143681
3	2.0	43525.0	44692.124842
4	2.2	39891.0	46582.117306
5	2.9	56642.0	53197.090931
6	3.0	60150.0	54142.087163
7	3.2	54445.0	56032.079627
8	3.2	64445.0	56032.079627
9	3.7	57189.0	60757.060788
10	3.9	63218.0	62647.053252
11	4.0	55794.0	63592.049484
12	4.0	56957.0	63592.049484
13	4.1	57081.0	64537.045717
14	4.5	61111.0	68317.030645
15	4.9	67938.0	72097.015574
16	5.1	66029.0	73987.008038
17	5.3	83088.0	75877.000502
18	5.9	81363.0	81546.977895
19	6.0	93940.0	82491.974127
20	6.8	91738.0	90051.943985
21	7.1	98273.0	92886.932681
22	7.9	101302.0	100446.902538
23	8.2	113812.0	103281.891235
24	8.7	109431.0	108006.872395
25	9.0	105582.0	110841.861092
26	9.5	116969.0	115566.842252
27	9.6	112635.0	116511.838485
28	10.3	122391.0	123126.812110
29	10.5	121872.0	125016.804574

In []: